

Introduction to Data Science Immersive (DSI)

Benjamin S. Skrainka

July 13, 2015

Today's schedule:

8:00 am:	Breakfast
8:30 am:	Welcome to Galvanize & introductions
9:00 am:	Galvanize community & tour
9:30 am:	Assessment & git
12:00 pm:	Lunch
1:00 pm:	Lecture (Ben)
2:15 pm:	Programming exercise 1 & portraits (Giovanna)

Welcome to Galvanize

Instructor introductions

Instructors for Seattle Cohort #1:

- ▶ Zach Alexander: zachary.alexander@galvanize.com
- ▶ Benjamin S. Skrainka: skrainka@galvanize.com
- ▶ Giovanna Thron: giovanna@galvanize.com

You can also reach us via Slack

Student introductions

Campus tour

Course expectations

Overview of course expectations

We will discuss basic logistics for DSI:

- ▶ Daily schedule
- ▶ Code of conduct
- ▶ Campus issues
- ▶ Logistics
- ▶ Assessment

Typical daily schedule

On most days, class will be a mix of lecture and exercises:

8:30 am:	Mini-quiz (survey on Fridays)
9:00 am:	Lecture
10:00 am:	Morning exercise (individual)
12:00 pm:	Lunch
1:15 pm:	Lecture
2:15 pm:	Afternoon exercise (pair)

Code of conduct

Campus issues

Logistics

DSI course uses `git` and [GitHub](#) to manage class materials and exercises.

- ▶ `git` is one of the best version control tools
- ▶ Fork repo(s) for each lecture and clone to your local machine
- ▶ Start with the Welcome repository
- ▶ Note: for legacy reasons the repo is owned by Zipfian

```
$ git clone git@github.com:joe_student/welcome.git
```

or

```
$ git clone https://github.com/joe_student/welcome.git
```

Assessment

Assessment consists of:

- ▶ Mini quizzes (most mornings)
- ▶ Individual exercise – usually in `individual.md`
- ▶ Pair programming exercise – usually in `pair.md`
- ▶ Roughly weekly assessment – usually in `assessment.md`

Regular feedback helps you track your progress and understanding of the material so that you can focus on the skills you need to improve

Lighting fast review of git

What is git?

Git is one of the best version control tools:

- ▶ Collaborate with many users around the Internet
- ▶ Keep changes synchronized across machines
- ▶ Can manage source, documentation, text files, and more
- ▶ A safety net:
 - ▶ Roll back mistakes
 - ▶ Explore new ideas safely (in a branch)
- ▶ Backup work!
- ▶ Increasingly, the standard tool for version control
- ▶ Reproducible research
- ▶ Unprofessional not to use git...

Git features

Features that make git awesome:

- ▶ SHA-1 hash:
 - ▶ Nigh uniquely identifies a commit across all copies of repo
 - ▶ Provides security from tampering
- ▶ Blazingly fast merges and diffs because git works on entire file and not incremental diffs
- ▶ Relatively painless merges
- ▶ Peer to peer:
 - ▶ Can work without Internet
 - ▶ Distributed
 - ▶ Fault tolerant

References

Two great references:

- ▶ [Pro Git](#)
- ▶ [git cheatsheet](#)

Plus, old-school help:

```
$ man git
```

```
$ git clone --help
```

Components of git

In order to understand git, you must understand how information flows between git's components:

- ▶ *Workspace*: current version of your work
- ▶ *Stage*:
 - ▶ Temporary container for work before committing to repository
 - ▶ Also known as the *Index* and *Cache*
- ▶ *Local repository*:
 - ▶ Local storage of all versions of your work which has been committed
 - ▶ Stored in `.git` directory in root of repo
- ▶ *Upstream repository*: remote copy, possibly out of sync with local repo
- ▶ *Stash*: 'park' work here when randomized by your boss

See [git cheatsheet](#)

Install git

Check if git is installed:

```
$ which git  
/usr/local/bin/git
```

If not, install [Homebrew](#) package manager and git:

```
$ ruby -e "$(curl -fsSL https://raw.githubusercontent.com/Homebrew/homebrew-core/master/install.sh)"  
$ brew install git
```

Create an account on GitHub

Setup up an account on [GitHub](#):

- ▶ To access course materials
- ▶ To submit your work
- ▶ To collaborate with students and colleagues
- ▶ To coordinate access to your repos on multiple machines
- ▶ To host your portfolio
- ▶ To backup your work

Configure git

Configure ~/.gitconfig and ~/.gitignore_global for your preferences. At a minimum, set the following:

```
$ git config --global user.name "Eddy Merckx"  
$ git config --global user.email eddy@merckx.com  
$ git config --global core.editor vim  
$ git config --global core.excludesfile ~/.gitignore_global
```

Eddy is a hardman, but you might prefer a friendlier editor, such as [Atom](#), nano, or pico

ONLY DO THIS ON YOUR OWN MACHINE

git clone *repo*

Create a local copy of the upstream repo:

```
$ git clone git@github.com:joe_student/awesome.git
```

Git workflow

Stage changes as you work and commit them to the local repo when you are ready:

```
$ atom genius.py  # Implement your great ideas
$ git status      # Check state of workspace and cache
$ git add genius.py  # Stage genius.py to commit to repo
$ git status
$ git commit -m 'Implemented orbital mind control laser'
```

Caveats

Do not put the following under version control:

- ▶ Large files: data, images, binary, .doc, .xls, .pdf, ...
- ▶ Derived files: i.e., files built from source code, markdown, \LaTeX , ...

Commit message should:

- ▶ Document why you did what you did
- ▶ May want to reference JIRA item or bug
- ▶ No need to say what you did
- ▶ Be nice to your future self!

Examining changes

You have several tools to compare versions:

- ▶ Use `tig` (Install via `brew install tig`)
- ▶ `git diff 16d6758 HEAD^^^`
- ▶ `git diff master`
- ▶ `git diff --stat`
- ▶ `git log`: see manual for details
- ▶ `git blame fubar.py` to determine who broke `fubar.py`

Syncing with an upstream repo (1/2)

To copy upstream changes to a local repo:

- ▶ Use `git pull`:
 1. Performs `git fetch` to create a local copy
 2. Attempts to merge changes via `git merge`
- ▶ Will require manual intervention if there is a merge conflict
- ▶ Can optionally specify the repo and branch to pull

Syncing with an upstream repo (2/2)

Use `git push` to copy local changes to an upstream repository:

- ▶ Must merge upstream changes before pushing your changes!
- ▶ May need to set an upstream tracking branch

```
$ git remote -v
origin  git@github.com:zipfian/bss.git (fetch)
origin  git@github.com:zipfian/bss.git (push)
$ git push
Counting objects: 15, done.
Delta compression using up to 4 threads.
Compressing objects: 100% (14/14), done.
Writing objects: 100% (15/15), 7.06 KiB | 0 bytes/s, done.
Total 15 (delta 3), reused 0 (delta 0)
To git@github.com:zipfian/bss.git
    16d6758..aac3438  master -> master
```

Submitting work via git

Typical class workflow:

1. Fork repo for lecture and exercise in GitHub
2. Clone your version of the repo onto local machine:

```
$ git clone https://github.com/joe/python-intro.git
```

3. Commit your work (do this often)

```
$ git add file1.py file2.py ...
```

```
$ git commit -m 'Implemented XYZ'
```

```
$ git push origin master
```

4. In GitHub,
 - 4.1 Click **Pull Requests**
 - 4.2 Click **New pull request**

Advanced commands

Some helpful commands:

- ▶ `git stash`: save changes in workspace so you can work on something else
- ▶ `git remote`: create and examine *remotes* to access upstream repositories
- ▶ `git init`: create a local repository in current directory
- ▶ `git reset`: rollback to previous commit
- ▶ `git revert`: revert commit(s)
- ▶ `git checkout`:
 - ▶ Revert file to version in local repo
 - ▶ Switch to a branch (create it if necessary)
- ▶ `git branch`: manipulate branches
- ▶ `git rebase`:
 - ▶ Replay changes from one branch on another
 - ▶ Smash small commits into a single commit

Pro tips

To improve your workflow:

- ▶ Use ssh and not https
- ▶ Use `tig`
- ▶ Work in a branch
- ▶ Commit and push regularly so that you can roll back changes if you make a mistake
- ▶ Never put keys in your repo; e.g., keys for AWS
- ▶ Be paranoid: enable dual-factor authentication
- ▶ Keep all repositories in `~/repos` or equivalent

Assessment

Using nose

Use nose to test your work:

- ▶ How to run tests
- ▶ How to interpret output

Lunch

Introduction to Data Science Immersive (DSI)

Objectives

Today's lecture focuses on:

- ▶ Setup your computer to do data science
- ▶ Install/configure core tools
- ▶ Adopt a productive workflow, both for data science and programming
- ▶ Review python Essentials
 - ▶ How to write efficient code
 - ▶ How to ensure your code is correct

Will discuss OOP tomorrow...

Agenda

1. Data science software stack

1.1 Overview data science and programming workflows

1.2 Introduction to core tools

2. Essential Python

2.1 Writing efficient code

2.2 $O(n)$ solution to anagrams problem

2.3 Looping efficiently

2.4 Using sets & dictionaries

Data Science Software Stack

Overview of key tools

We focus on the key tools you need to work on data science projects:

- ▶ ipython
- ▶ iTerm2
- ▶ Atom or other suitable programming editor (vi, emacs, Sublime Text)
- ▶ Unix

Plus, an overview of typical workflow

Why Python for data science?

Python is as close to an all purpose language as you can find:

- ▶ Strong for programming and integration tasks
- ▶ Interpreted \Rightarrow easy to prototype and explore
- ▶ Supports most key data science technologies
- ▶ Excellent machine learning syntax
- ▶ Simple syntax

But, R often has superior statistics packages . . .

References

A couple helpful references on Python and the craft of writing software:

- ▶ A Practical Guide to Commands, Editors, and Shell Programming
- ▶ Effective Computation in Physics
- ▶ Python for Data Analysis
- ▶ “Effective Python”
- ▶ Skrainka’s [lecture](#) on writing software from ICE
- ▶ “The Practice of Programming”
- ▶ “Code Complete”

Installation

To install the data science stack used at Galvanize:

- ▶ Do nothing, if on a lab machine
- ▶ On your laptop, run commands in `setup.sh`:
 - ▶ Contained in *Welcome* repo
(`git@github.com:zipfian/welcome.git`)
 - ▶ Run commands by hand because many may time-out and/or require baby-sitting
- ▶ Installation uses [Homebrew](#), a package manager for OS/X
 - ▶ Brew is incompatible with [MacPorts](#)
 - ▶ On Linux, use the package manager, e.g., `apt-get` on Ubuntu

Invest in Unix

Unix's *filters + pipes* paradigm provides a productive platform for doing data science:

- ▶ Each command does one thing (only) well
- ▶ Commands read from `STD_IN` and write to `STD_OUT`
- ▶ Can combine commands via I/O redirection with `|`, `<`, and `>`.
- ▶ Easy to write scripts to build bespoke commands from Unix's building blocks

Navigating in Unix

Basic survival commands for Unix:

Command	Action
<code>pwd</code>	Display current directory
<code>mkdir</code>	Create a directory (folder)
<code>ls</code>	Display contents of folder
<code>cd</code>	Change directory
<code>man</code>	Get help (manual)

Finding stuff

Unix provides several great tools for finding stuff:

Command	Action
<code>ack</code>	recursively search files in a tree for an RE pattern
<code>grep</code>	original tool for searching for a pattern in a file
<code>find</code>	recursively perform operations on directory tree
<code>locate</code>	find a file on computer
<code>which</code>	determine which version of command is accessible

Hacking on text

Unix also provides great tools for processing text files:

Command	Action
<code>file</code>	Display file type
<code>wc</code>	Count lines, words, and characters
<code>od</code>	Dump file in octal
<code>sort</code>	Sort file
<code>uniq</code>	Display unique lines or histogram
<code>cut</code>	Cut columns from text file
<code>paste</code>	Paste two text files
<code>sed</code>	Stream editor
<code>awk</code>	Programmable stream editor

and more!

On OS/X, use iTerm2

iTerm2 is a better version of the Terminal app:

- ▶ Provides command line access to bash or other shell
- ▶ Decreased eye strain
- ▶ Better support for multiple panes and hotkeys
- ▶ Improved search facility with highlighting
- ▶ Mouse-less copy
- ▶ Paste history
- ▶ More configurable

```
$ brew cask install iterm
```


Overview of data science workflow

A data science project typically follows a model similar to OSEMN:

- ▶ **O**btain the data
- ▶ **S**crub the data
- ▶ **E**xplore the data
- ▶ Build a **M**odel
- ▶ **I**Nterpret the results

May need to iterate over these steps.

“CRISP-DM Process Diagram” by Kenneth Jensen - Own work. Licensed under CC BY-SA 3.0 via Wikimedia Commons



Reproducible research

Make sure your work is reproducible:

- ▶ Often, DS projects are repeated (seasonally)
- ▶ Often, someone will have to take over a project from you
- ▶ Often, you will have to resume a project after a lengthy break
- ▶ Consequently, make sure you keep a (lab) notebook and document:
 - ▶ Data provenance
 - ▶ Data cleaning
 - ▶ Model and assumptions
 - ▶ Data science pipeline
- ▶ Use git to manage code and documentation

Overview of tactical workflow

To be productive, you want a tight feedback loop:

- ▶ Write some code
- ▶ Run it/test it to get quick feedback
 - ▶ Easier to fix a bug if you catch it early
 - ▶ Maintenance is the primary cost of building software
- ▶ Start small, and get the simplest case to work first
- ▶ If working with data, get a shard (subset) to work first
 - ▶ Will provide quick feedback that your code works
 - ▶ Build a shitty first model, then refine
 - ▶ Do not start with an Exabyte of data!

Become a better programmer

Invest in becoming a better programmer:

- ▶ Increased productivity
- ▶ Ability to solve bigger and harder problems
- ▶ Master the craft as well as grammar:
 - ▶ Write more code
 - ▶ Get code reviews
 - ▶ Read well-written code
- ▶ Try to use the best tool for the job:
 - ▶ Learn more languages
 - ▶ R, Scala, Julia, Java, MATLAB, Mathematica, Clojure,...

Working in Python

Depending what you are doing you will use one of the following workflows, often in conjunction with *Test Driven Development* (TDD):

- ▶ Editor + python or ipython
- ▶ ipython notebook

TDD emphasizes working with unit tests to ensure a tight feedback loop and avoid the debugger (see below)

Invest in learning an editor

Master a programming editor, it will save you immense amounts of time:

- ▶ Color-syntax highlighting and line numbers
- ▶ Can configure PEP8 syntax checking to catch mistakes early (when cheaper to fix)
- ▶ Automates much of writing code
- ▶ Fancy search and text manipulation
- ▶ Quickly navigate a large code base
- ▶ Use [Atom](#) unless you have already mastered vi(m) or emacs

Working in ipython

ipython makes Python more friendly for data science:

- ▶ %magic methods plus `ls`, `cd`, `pwd`, etc. to navigate file-system
- ▶ Tab completion
- ▶ Enhanced help and introspection via `?`
- ▶ Keyboard shortcuts
- ▶ Convenience variables:
 - ▶ Access past results with `_`, `__`, `_72`, etc.
 - ▶ Access past inputs with `_i` and number of command
 - ▶ Can re-execute command with `exec _i99`
- ▶ History
- ▶ Integrated plotting

A little %magic

ipython offers many 'magic' methods to improve productivity:

- ▶ `%run file.py`: execute a script
- ▶ `%debug`: enter debugger at source of last exception
- ▶ `%pdb`: automatically enter debugger on exception
- ▶ `%timeit func`: compute average execution time of *func*
- ▶ `%paste` and `%cpaste`: execute code from clipboard or pasting
- ▶ `%reset`: remove everything from namespace

Working in ipython notebook

But, ipython notebook is even better:

- ▶ Provides notebook to document work as well as run python commands
 - ▶ Similar to Mathematica . . .
 - ▶ Promotes reproducible research by unifying code, documentation, and results
 - ▶ Inline plotting
- ▶ Use for exploratory data analysis (EDA) and prototyping
- ▶ Do not become sloppy:
 - ▶ Write functions, classes, modules, etc. for production code
 - ▶ . . . or to reuse common code
 - ▶ Prefer DRY to WET coding
- ▶ Can also run on remote machine and connect via ssh tunnel
- ▶ Covered in Thursday's lecture

Python Essentials

Review of Python Essentials

Objectives for the rest of lecture:

1. Writing efficient code
2. $O(n)$ solution to anagrams problem
3. Looping efficiently
4. Using sets & dictionaries

Basic software design

A few basic, old school tips:

- ▶ Good programmers are lazy
- ▶ Start by listing requirements and/or writing a specification
- ▶ Simplify complex problems by 'divide and conquer'
- ▶ Decompose problem into functions:
 - ▶ An interface is a contract
 - ▶ Each function does one thing only and does it well
 - ▶ If your code spans a screen or two, break it up!
- ▶ Structured programming / Top down / Bottom up

More basic design

Algorithms and data structures complement each other:

- ▶ Choose the right data structure for the job:
 - ▶ E.g., what is the correct container (list, dict) to use?
- ▶ Reuse code via libraries wherever possible:
 - ▶ Code is already debugged
 - ▶ Interfaces are battle-tested
- ▶ Does it make sense to sort?
- ▶ For only a few elements, brute force is fastest
- ▶ Fail early

DRY vs. WET

Be DRY not WET:

- ▶ DRY means 'Do not Repeat Yourself'
- ▶ WET means 'We Enjoy Typing'

If you are writing code with cut & paste, you really need to write a function:

- ▶ Promotes reuse
- ▶ Decreases bugs
- ▶ Improves maintainability

Don't let ipython notebook turn you into a sloppy programmer!

Scope: LEGB

Questions of scope are resolved using LEGB:

1. Locals
2. Enclosing
3. Globals
4. Built-ins

Deep vs. shallow copy

Python passes object references by value:

- ▶ Like passing a pointer to an object in C:
`status = serialize_data(&data) ;`
- ▶ Python uses object reference to refer to an object
 - ▶ Saves memory
 - ▶ Can accidentally modify objects because default copy is shallow
- ▶ deep-copy objects when you want a copy of the object and any objects it contains
- ▶
- ▶ Can see object's ID with `id(obj)`

See: [Rob Heaton's post](#)

Example: deep vs. shallow copy

```
In [1]: x = [ [1, 2, 3], [-4, -5, -6]]
```

```
In [2]: [id(ix) for ix in x]
```

```
Out[2]: [4419005976, 4419004536]
```

```
In [3]: y = x
```

```
In [4]: [id(ix) for ix in y]
```

```
Out[4]: [4419005976, 4419004536]
```

```
In [5]: y[1][0] = 'four'
```

```
In [6]: y
```

```
Out[6]: [[1, 2, 3], ['four', -5, -6]]
```

```
In [7]: x
```

```
Out[7]: [[1, 2, 3], ['four', -5, -6]]
```


Mutable vs. immutable

Data type can be:

- ▶ Mutable:
 - ▶ list, set, dict
 - ▶ Can be modified after creation
- ▶ Immutable:
 - ▶ str, int, float, tuple
 - ▶ Cannot be modified after creation

Looping efficiently

To write faster loops:

- ▶ Prefer `for word in words:` to direct indexing with `for ix in xrange(len(words)):`
- ▶ Prefer `enumerate()` to `xrange()` when you need the index and the value
- ▶ Prefer `for k, v in my_dict.iteritems():` to `for k in my_dict:` when you need both the key and value
- ▶ Prefer `itertools.izip()` to indexing over multiple lists
- ▶ Prefer `itertools.izip()` to `zip()`

Sets & Dictionaries

Sets and dictionaries provide fast look-up, insertion, and deletion:

- ▶ Prefer `for k, v in my_dict.items():` to `for k in my_dict.keys():` if you need both the key and value
- ▶ Check membership with `in`
- ▶ Prefer `Counter` and `defaultdict` from `collections`
 - ▶ If order matters, prefer `OrderedDict` or `OrderedSet`
- ▶ Implemented as hash table so should have roughly $O(1)$ access speed
- ▶ Note: dictionary keys and set items must be immutable

Solving anagrams in $O(n)$

Recall the interview question: *write a function to return all the words which have anagrams in a list of words*

- ▶ Q: what is the algorithmic cost of the double for loop solution?
- ▶ Q: can you think of an algorithm which is $O(n)$?
 - ▶ Take a few minutes to ponder and discuss with a partner
 - ▶ Hint: what data structure has fast, random access to elements?
 - ▶ Hint: what does defaultdict do?
 - ▶ Hint: how can you transform each word so that it is easy to find an anagram?

Anagram timing

How much better is the $O(n)$ solution?

```
In [18]: %timeit anagrams.anagrams1(anagrams.words)
10000 loops, best of 3: 46.2 µs per loop
```

```
In [19]: %timeit anagrams.anagrams2(anagrams.words)
100000 loops, best of 3: 10.4 µs per loop
```

List comprehensions (1/2)

List comprehensions provide a compact way to create a list similar to a for loop

- ▶ Enclose the for expression in []

- ▶ Returns a list with elements:

```
squares = [val**2 for val in xrange(10)]
```

- ▶ List comprehension is equivalent to:

```
squares = []  
for val in xrange(10):  
    squares.append(val**2)
```

List comprehensions (1/2)

List comprehensions are flexible and can be customized:

- ▶ Nested for loops:

```
L = [[1, 2, 3], [4, 5, 6], [7, 8, 9]]  
flat_list = [item for row in L for item in row]
```

- ▶ Add a condition with if:

```
even_squares = [  
    val**2 for val in xrange(10) if 0 == val % 2]
```

Generators

Generators are a superior form of an 'iterable' which provide values on demand:

- ▶ Decrease memory usage via lazy evaluation
- ▶ Often used in a loop, e.g., `xrange()`
- ▶ Increase performance with large data because they generate the next value on each iteration:
 - ▶ `range(100000)` will create the entire list of values at once
 - ▶ `xrange(100000)` will create the next value on each iteration
- ▶ Prefer generators over classic iterables
- ▶ Can create via generator comprehension or with `yield`

See: [David Beazley's talk](#)

Generator comprehensions

Generator comprehensions resemble list comprehensions and provide a quick way to create an iterator:

- ▶ Replace the `[]` with `()` to create a generator
- ▶ Will only produce the next value on demand, saving memory
- ▶ No need to worry about implementation details such as `.next()`, `.__iter__()`, and `yield`
- ▶ But, generator is consumed when you use it, so one use only
- ▶ Benefits:
 - ▶ Faster
 - ▶ Uses less memory
 - ▶ Requires less code

Generator example (From David Beazley)

```
wwwlog = open("my_web_log")
bytecol = (line.rsplit(None,1)[1] for line in wwwlog)
bytes    = (int(x) for x in bytecol if x != '-')
print "Total:", sum(bytes)
```

Lambda

Provide a handy way to write simple, anonymous functions:

- ▶ Useful for simple reductions and aggregations

- ▶ Syntax:

- ▶ `lambda <args>: <body>`
- ▶ Can have multiple arguments
- ▶ Do not include a return statement
- ▶ Should be simple, otherwise write a function

- ▶ Lambdas are often used with `map()` and `reduce()`

- ▶ Example:

```
df = pd.read_csv('awesome_data.csv')  
df.apply(lambda x: x.max() - x.min())
```

- ▶ Prefer `operator.itemgetter()` for extracting an element from a list/tuple to `lambda + indexing`

Using exceptions

Python (and your code) raises an exception to signal an exceptional event such as an error:

- ▶ Only catch errors you must handle, such as file I/O problems
- ▶ Let the error unwind the call stack for everything else:
 - ▶ Easier to read code
 - ▶ Forces user to fix exceptional condition
 - ▶ Improves performance of code
- ▶ Keep try blocks small
- ▶ Catch only the specific errors you need to handle

Context management with with

Context managers ensure that you cleanly allocate and deallocate a resource:

- ▶ Python handles setup and tear-down behind the scenes by calling object's `__enter__()` and `__exit__()` methods
- ▶ Ensures cleanup on an exception
- ▶ Often used with file I/O, or when obtaining a lock
- ▶ Example:

```
with open('logfile.txt', 'r') as f:  
    lines = f.readlines()
```

PEP8

PEP8 is the official recommendation for best Python practices:

- ▶ Always check your code with `pep8` to help catch bugs early
- ▶ Can run from command line
- ▶ Better to configure editor to automatically check code
 - ▶ For vim, install [Python-mode](#)
 - ▶ For Sublime, install *Package Control*, *PEP8 AutoFormat*, and *SublimeLinter*
 - ▶ For Atom, PEP8 support is already installed